

DESCRIBE THE ROLE AND RESPONSIBILITIES OF AN OS

The operating system is the most important program that runs on a computer. Every general-purpose computer must have an OS to run other programs. Operating systems perform basic tasks, such as recognizing input from the keyboard, sending output to the display screen, keeping track of files and directories on the disk, and controlling peripheral devices such as disk drives and printers. The operating system acts as an intermediary between a user of a computer (and the used application programs) and the computer hardware. It also provides environment for other software to execute correctly.

DESCRIBE THE GENERAL ORGANIZATION OF A COMPUTER SYSTEM AND THE ROLE OF INTERRUPTS.

The computer system can be divided into four components:

THE HARDWARE

- CPU
- Memory
- I/O devices
- Provides basic computing resources for the system.

THE OPERATING SYSTEM

Controls the hardware and coordinates its use among the various application programs for the various users. The operating system acts as a resource manager.

THE APPLICATION PROGRAMS

- such as word processors,
- spreadsheets,
- compilers,
- Web browsers

THE USER

Communicates with the system through the user interface, such as the command-line interpreter or graphical user interface (GUI).

ROLE OF INTERRUPTS

Interrupt service routines (ISR) preserves the state of the CPU by saving registers and then calls the appropriate OS routine to handle the interrupt. When the OS routine completes, control is returned to the ISR, which restores the saved registers and returns control to the interrupted program.

DESCRIBE THE COMPONENTS IN A MODERN MULTIPROCESSOR COMPUTER SYSTEM.

A multiprocessor system consists of **TWO** or more CPUs that share a common physical memory. Multiprocessor systems are also known as **parallel systems, tightly coupled systems, and shared-memory systems.**

Multiprocessor systems are more complex than uniprocessor systems because of the **need to manage concurrent access to the shared memory.** These systems are more economical because they can share resources.

Multiprocessor systems can be categorized according to the number of CPUs. Symmetric multiprocessing systems (SMP) and asymmetric multiprocessing systems (AMP).

- SMP's have two or more similar processors running the same OS and performing the same tasks.
- AMP's have one master processor and one or more slave processors.
 - The master processor schedules and allocates work to the slave processors.
 - A clustered system consists of two or more individual systems joined together.

- The individual systems are independent but work together as a single system.

ILLUSTRATE THE TRANSITION FROM USER MODE TO KERNEL MODE.

The transition from user mode to kernel mode occurs when a user program requests a service from the OS, such as a request to read data from a file. The system must ensure that the request is valid and that the user program has the right to access the file. The system then executes the request on behalf of the user program.

The transition from user mode to kernel mode is usually done via a system call, which is a request to the OS to allow a user program to access a resource. The system call is usually initiated by a user program via a software interrupt. The system call is handled by a dispatcher, which is a routine within the OS that examines the request and determines how to execute it.

The dispatcher then invokes the appropriate OS routine to perform the request. When the OS routine completes, control is returned to the dispatcher, which returns control to the user program.

HOW ARE OS'S USED IN VARIOUS COMPUTING ENVIRONMENTS.

Operating systems are used in a variety of computing environments, including:

Desktop-, multiprocessor-, distributed-, cluster-, real-time- and handheld systems

- **Multiprocessor** systems are used to increase throughput and reliability.
- **Distributed** systems are used to provide users with access to remote resources, such as printers, files, and databases.
- **Cluster** systems are used to provide high availability and load balancing.
- **Real-time** systems are used as control devices in a dedicated application.
- **Handheld** systems are used to provide computing resources in a small, portable package.

PROVIDE EXAMPLES OF FREE AND OPEN-SOURCE OPERATING SYSTEMS.

Examples of free and open-source operating systems include:

Linux

FreeBSD

NetBSD

OpenBSD

Linux is a free and open-source OS that is based on UNIX and is available for a wide range of computing platforms. Linux is available in a number of distributions, such as Red Hat, Fedora, Ubuntu, and Debian.

FreeBSD, **NetBSD**, and **OpenBSD** are free and open-source OSs that are based on UNIX and are available for a wide range of computing platforms.

Vika 4

IDENTIFY SERVICES PROVIDED BY AN OPERATING SYSTEM.

- **Error detection:** Detect errors in the CPU and memory hardware
- **Program execution:** Load a program into memory and run it
- **I/O operations:** Transfer data to and from I/O devices
- **File-system manipulation:** Read, write, create, delete and search files and directories
- **Communications:** Exchange information between processes executing either on the same computer or on different systems tied together by a network
- **Resource allocation:** Allocate resources to multiple users or multiple jobs running at the same time

- **Accounting:** Keep track of which users use how much and what kinds of computer resources
- **Protection and security:** Protect the computer and its data from unauthorized use, either by ensuring that only authorized users are allowed access to the system or by protecting individual files and other system resources against unauthorized access. It also makes sure that concurrent processes running in the system do not interfere with each other.

HOW ARE SYSTEM CALLS USED TO PROVIDE OS SERVICES.

The system call API invokes intended system call in the kernel, by passing number (and additional parameters) using a trap assembly instruction, which then performs the requested service and returns control to the caller.

COMPARE AND CONTRAST MONOLITHIC, LAYERED, MICROKERNEL, MODULAR, AND HYBRID STRATEGIES FOR DESIGNING OS'S.

Monolithic: The kernel is a single program that provides all of the services of the operating system. It executes in kernel mode and has access to all of the hardware and data structures of the kernel. It is a single static binary file.

Layered: The operating system is divided into a number of layers (levels), each built on top of lower layers. The bottom layer (layer 0), is the hardware; the highest (layer N) is the user interface. With modularity, layers are selected such that each uses functions (operations) and services of only lower-level layers. If a layer is changed, the layers above it are unaffected. The major difficulty with the layered approach involves defining the layers.

Microkernel: Moves as much from the kernel into user space. Communication takes place between user modules using message passing. Benefits include: easier to extend a microkernel; easier to port the operating system to new architectures; more reliable (less code is running in kernel mode); more secure (less code is running in kernel mode). Disadvantages include: performance overhead of user space to kernel space communication; increased size of operating system.

Modular: Instead of having a single, monolithic kernel, the kernel is broken down into separate processes, known as servers. Some of the servers are:

- file server,
- process server,
- and memory server.

The servers invoke system calls as needed by sending messages to other servers. The kernel is not a separate entity, but is a set of cooperating processes in user space.

Benefits include:

- easier to extend a modular operating system;
- easier to port the operating system to new architectures;
- more reliable (less code is running in kernel mode);
- more secure (less code is running in kernel mode).

Disadvantages include:

- performance overhead of user space to kernel space communication;
- increased size of operating system.

Hybrid: Combines the speed of a microkernel with the modularity of a modular kernel. The kernel consists of a microkernel, but the servers are divided into modules, each running in user space. The microkernel provides minimal process and memory management, interprocess communication, and basic synchronization primitives. The kernel modules provide the file system, device drivers,

networking, and other operating system functions. The kernel modules can be loaded and unloaded dynamically, making it easier to extend the kernel.

ILLUSTRATE THE PROCESS FOR BOOTING AN OPERATING SYSTEM.

- The BIOS is located in ROM on the motherboard. It is the first code that is executed at start-up and is responsible for locating and loading the operating system kernel software.
- The BIOS performs a power-on self-test (POST) to ensure that all of the hardware components are present and operational.
- The BIOS then loads the first sector of the boot device (usually a hard disk) into memory and transfers control to that code.
- This code is known as the master boot record (MBR).
- The MBR locates the active partition on the hard disk and loads a copy of its first sector into memory.
- This code is known as the volume boot record (VBR).
- The VBR loads the operating system kernel into memory and transfers control to it.
- The kernel initializes the rest of the operating system.
- The kernel creates a process for the init program, which is the first user-level process.
- The init program then starts other processes, such as daemons, which are background processes that provide services to the system.
- The init program waits for the system to shut down or reboots the system if instructed to do so.

Vika 5

IDENTIFY THE SEPARATE COMPONENTS OF A PROCESS AND ILLUSTRATE HOW THEY ARE REPRESENTED AND SCHEDULED IN AN OPERATING SYSTEM.

Program counter (PC): The PC is a register that contains the address of the next instruction to be executed.

Stack: The stack is a data structure that contains temporary data such as function parameters, return addresses, and local variables.

Data section: The data section contains global variables.

Set of further associated resources like **heap** and **open files**

DESCRIBE HOW PROCESSES ARE CREATED AND TERMINATED IN AN OPERATING SYSTEM, INCLUDING DEVELOPING PROGRAMS USING POSIX SYSTEM CALLS THAT PERFORM THESE OPERATIONS.

Process creation:

- The **fork()** system call creates a new process by **duplicating** the calling process. The new process is referred to as the **child** process. The calling process is referred to as the parent process.
- The **exec()** system call used after a fork to replace the process' memory space with new program.

Process termination

- **wait()** returns data from child to parent (return value provided by exit system call)
- **exit()** process executes last statement and voluntarily requests from the OS to be deleted.

DESCRIBE AND CONTRAST INTERPROCESS COMMUNICATION USING SHARED MEMORY AND MESSAGE PASSING.

Shared memory:

- Once the shared memory has been established, ordinary memory access techniques can be used to exchange information between processes without further OS support.
- Processes need to synchronize their access to the shared memory to avoid conflicts.
- Preferable when large amounts of data need to be exchanged between processes.

Message passing:

- OS has an internal buffer that can be accessed by different processes via send and receive operations to exchange data.
- Operating system provides a set of system calls to create and manage the message buffers and to send and receive messages.
- Preferable when smaller amounts of data need to be exchanged between processes.

DESCRIBE PROGRAMS THAT USE POSIX PIPES AND POSIX SHARED MEMORY TO PERFORM INTERPROCESS COMMUNICATION.

Pipe: Pipes are a form of IPC that allow data to be transmitted between processes in a linear, unidirectional manner. A pipe consists of a read end and a write end. When one process writes data to the write end of the pipe, another process can read that data from the read end. Pipes can be used to implement filters, where the output of one process serves as input for another process.

Shared memory: Shared memory is a region of memory that can be accessed by multiple processes simultaneously. It's an efficient form of IPC because it doesn't require copying data between processes; instead, they read and write directly to the shared memory region. Synchronization primitives like semaphores or mutexes are often used to ensure the integrity of the data in shared memory.

DESCRIBE CLIENT-SERVER COMMUNICATION USING SOCKETS AND INCLUDING HOW TO CREATE CLIENT/SERVER PROGRAMS USING THE JAVA SOCKET API.

Sockets: A socket is one endpoint of a two-way communication link between two programs running on the network. A socket is bound to a port number so that the TCP layer can identify the application that data is destined to be sent to. An endpoint is a combination of an IP address and a port number. Every TCP connection can be uniquely identified by its two endpoints. That way you can have multiple connections between your host and the server.

Java sockets: Java API supports interprocess communication using sockets. The `java.net` package provides classes that represent sockets and server sockets. The Port numbers are represented by `Integers`. IP addresses are represented by `InetAddress` objects. And for translating domain names to IP addresses, the `InetAddress` class provides a static method called `getByName()`.

IDENTIFY THE BASIC COMPONENTS OF A THREAD, AND CONTRAST THREADS AND PROCESSES.

A thread is a basic unit of CPU utilization; it comprises a thread ID, a program counter, a register set, and a stack. Creation of a thread and context switch is more efficient/faster than that of a process. Threads share the same address space, while processes have their own address space. Threads are used for concurrency, while processes are used for (full) parallelism. Threads are more lightweight than processes.

DESCRIBE THE MAJOR BENEFITS AND SIGNIFICANT CHALLENGES OF DESIGNING MULTITHREADED PROCESSES.

Benefits:

- **Responsiveness**, a multithreaded process may start one thread for computation, one thread for user interaction, etc
- **Resource sharing**, memory of process is shared between threads, no system calls required for creating shared memory area or for message passing
- **Economy**, creating threads is faster than creating processes. Context switch (between threads of same process) is faster for threads than for processes
- **Scalability**, each thread can be executed by a different processor/core, achieving a speed-up by parallel processing.

Challenges:

- **Dividing activities**, which activities can run in parallel
- **Balance**, overhead of thread handling / communication / synchronisation may outweigh performance gain
- **Data splitting**, not only a challenge how to split activities, but also how to divide data processed by different threads
- **Data dependency**, if a thread depends on data produced by another thread, synchronisation between threads is needed
- **Testing and debugging**, inherently more difficult than single-threaded applications.

DESCRIBE DIFFERENT MULTITHREADING MODELS.

Many-to-one model: Many user-level threads mapped to one kernel thread.

One-to-one model: One user-level thread mapped to one kernel thread.

Many-to-many model: Many user-level threads mapped to many kernel threads.

Two-level model: A combination of the many-to-one and one-to-one models.

DESIGN MULTITHREADED APPLICATIONS USING THE POSIX PTHREADS AND JAVA THREADING APIs.

1. create new thread and return its thread Id:

```
pthread_create(  
    pthread_t *thread,  
    const pthread_attr_t *attr,  
    void *(*start_routine) (void *),  
    void *arg  
);
```

2. Thread hands over control to thread library:

```
int pthread_yield()
// or
int sched_yield()
```

3. Thread terminates itself:

```
void pthread_exit(void *retval)
```

4. Thread waits for termination of another thread:

```
int pthread_join(pthread_t thread, void **thread_return)
```

5. Java thread creation:

```
class MyThread extends Thread {
    public void run() {
        // code to be executed in new thread
    }
}

class MainThread {
    public static void main(String args[]) {
        MyThread t = new MyThread();
        t.start();
    }
}
```

HAVE HEARD ABOUT IMPLICIT THREADING APPROACHES.

Implicit threading: Threads are created and managed by compilers and runtime libraries instead of programmer.

Advantages:

- Programmer does not need to worry about thread creation, management, and synchronization.
- Compiler and runtime library can decide how to map threads to processors.

Disadvantages:

- Programmer has less control over thread management.
- Compiler and runtime library may not be able to determine how to parallelize the code.

Vika 7

MOTIVATION FOR SCHEDULING

Only one process/thread can run on a processor (or core) at a time. All other processes must wait until CPU is free, and they are scheduled.

MULTIPROGRAMMING (BATCH) SYSTEM

Maximise CPU utilisation and throughput of jobs. While one process is blocked (i.e. due to I/O) another process may use the CPU.

TIMESHARING (MULTITASKING) SYSTEM

While one process/thread is performing calculations, user can still interact with another process/thread because scheduler switches often between them.

WHY SCHEDULING IS REASONABLE

Overhead of scheduling is generally outweighed by the benefits. Process execution typically consists of a period of CPU usage and subsequent I/O wait. During this wait scheduling enables another process to utilise the CPU.

DEFINITIONS

- **(CPU) Scheduler:** Part of the OS kernel that assigns CPU time to processes/threads that are ready to execute.
- **Dispatcher:** Part of the OS kernel that performs the actual context switch (restoring CPU registers, switching from kernel to user mode)
- **Scheduling algorithm:** The algorithm used by the scheduler to decide which process/thread gets the CPU for how long.

REMARKS

- There is no **best** scheduling algorithm. Different algorithms are best suited to different types of system (batch, multitasking, real-time, etc.) and usage scenarios.
- In operating systems with kernel level threads, **only** threads (not processes) are scheduled.

PREEMPTIVE VS NON-PREEMPTIVE SCHEDULING

NON-PREEMPTIVE

CPU is allocated to one process until that process blocks or terminates. Timesharing is only possible if CPU bound processes explicitly yield the CPU, by using the yield system call to voluntarily transition from **RUNNING** to **READY** state.

PREEMPTIVE

A running process may be interrupted at any time because its time slice has expired. The scheduler then takes control and determines which process gets to use the CPU next. This may be the same process, or any other process in **READY** state.

POTENTIAL PROBLEMS

- A process may be interrupted while updating data shared between processes. Cooperating processes must therefore synchronise access to such resources to ensure a single process has exclusive access to the data until it completes updating it.
- A time slice timer can expire while kernel code is being executed, the kernel must therefore disable interrupt processing while updating critical kernel data structures.

SCHEDULING CRITERIA

All scheduling systems must ensure:

- **Fairness:** Each process gets CPU time
- **Enforcement of priorities:** High priority processes are preferred.
- **Balance:** All the different resources of a system are reasonably utilised.

Batch system schedulers must also ensure:

- **CPU utilization:** The CPU should be kept as busy as possible.
- **Throughput:** The number of processes that complete their execution per time unit should be maximised.
- **Turnaround time:** Minimise the amount of time (from start to termination) to execute a particular process.

An interactive system must ensure:

- **Response time:** Minimise the amount of time it takes from the time a request is submitted until the first response is produced.

A real-time system must ensure:

- **Meeting deadlines:** Processes (or events within processes) that must be started/finished before a certain point in time must be preferred.
- **Predictability:** As long as the system is not overloaded, it can be predicted when a certain process (or event within) is executed.

SCHEDULING ALGORITHMS

FCFS, SJF, and SRTF are primarily applicable to Batch operating systems, since each process runs more or less to completion. To present the illusion of multiple processes running simultaneously, interactive operating systems must employ different algorithms, such as Round Robin and its variants.

FIRST COME FIRST SERVED (FCFS)

In a FCFS scheduler processes are allocated CPU time in the order of arrival, and running processes are not interrupted. This means that the first process to arrive will run to completion, before the second gets the CPU and runs to completion, etc. This algorithm is non-preemptive, easy to implement, and fair (in the sense that all processes will eventually get access to the CPU). For this reason the average waiting time for a process, and the general suitability of the algorithm are heavily dependent on the order in which processes are created.

SHORTEST JOB FIRST (SJF)

In a SJF scheduler processes are served in ascending order of CPU time required (based on the processes in queue at the time of scheduling decisions). This imposes the limitation that the CPU time required by a process must be known in advance (unlikely in real world scenarios). This algorithm suffers from the problem that it is unfair, since a process requiring large amounts of CPU time will never be executed if shorter processes keep arriving.

SHORTEST REMAINING TIME FIRST (SRTF)

The SRTF algorithm is a preemptive variant of SJF, where the process running process is interrupted if a newly arrived process requires less CPU time than the running process would require to complete. This algorithm still suffers from the unfairness problem of SJF, where a long process will never get the CPU if shorter processes keep arriving.

ROUND-ROBIN (RR)

In a Round-Robin scheduler CPU time is divided into **time slices** with a fixed maximum duration (If a process completes before its time slice expires, the next process does not wait for the time slice to expire, but starts immediately.) Processes are then served in a First Come First Served manner (with new processes simply placed at the back of the READY queue), with each process getting the CPU for one time slice, before being placed at the back of the queue. When a process terminates it is removed from the READY queue. If a process blocks, i.e. due to I/O it is removed from the READY queue and placed onto the WAITING queue until its blocking request has been satisfied, at which point it reenters the READY queue.

ROUND-ROBIN WITH PRIORITIES

In a Round-Robin with Priorities scheduler the READY queue is replaced by multiple queues, where each queue has a priority value. The highest priority queue is processed in a Round-Robin fashion,

and only once it is empty is the next queue processed. This has the potential to cause starvation in low priority processes, which can be countered by dynamically adjusting the priority of processes (Increase the priority of processes that have spent a long time waiting, decrease the priority of long running processes).

MULTILEVEL QUEUE SCHEDULING

Different categories of processes (interactive, background, system, etc.) are placed in different queues. Each queue has a different scheduling algorithm. Some sort of algorithm is required to choose which queue gets to run.

MULTILEVEL FEEDBACK QUEUE SCHEDULING

Multilevel Queue Scheduling, except processes can be moved between queues.

THREAD SCHEDULING

If user level threads are used, the OS kernel is not aware of the existence of the threads, but simply schedules the processes. Scheduling of the threads is left to the user level thread library. If kernel level threads are used the kernel schedules threads, and typically does not care to which process those threads belong.

MULTIPLE-PROCESSOR SCHEDULING

When more than one CPU/core are present in a system, and share memory, each core must be managed by the operating system.

ASYMMETRIC MULTIPROCESSING

Only one **master** processor/core accesses the kernel data structures (such as scheduler queues). Other processors (**slaves**) wait for the master processor to assign them work.

SYMMETRIC MULTIPROCESSING (SMP)

All processors/cores run the same kernel, and make independent scheduling decisions. This is the scheme used by all major operating systems these days. This can either be implemented by a shared scheduler queue, access to which must then be synchronised, or each processor can maintain its own scheduler queue.

PROCESSOR AFFINITY

Since each CPU core has its own cache for recent data and instructions it is inefficient to constantly move processes between cores, and thereby invalidate all caches, requiring costly memory accesses, the scheduler tries to keep a process on the same physical core. The process is then said to have affinity for that processor.

HARD PROCESSOR AFFINITY

In a hard affinity model a process is never moved between processors, such as when each processor has its own scheduler queue. Under this model some cores may sit idle, even though processes are waiting in queue, because they have an affinity for a different processor.

LOAD BALANCING

In opposition to Processor Affinity, load balancing attempts to evenly distribute workload between available processors.

PUSH MIGRATION

In a push migration scheme the kernel periodically checks the load on each processor and migrates (pushes) processes from cores with high load, onto cores with light load.

PULL MIGRATION

In a pull migration scheme a processor whose scheduling queue is empty will pull processes from another processor's queue.

SOFT PROCESSOR AFFINITY

Load Balancing and Processor Affinity contradict each other, and it is difficult to develop algorithms that achieve a good compromise between the two. Such attempts are known as soft processor affinity, and revolve around attempting to maintain affinity, but allowing load balancing where necessary.

MEMORY STALLS

On the OS level processes get blocked while waiting for things like I/O. The same may happen on the CPU level, since main memory access is significantly slower than the CPU itself. When this occurs it is known as a memory stall, and results in wasted CPU cycles. Memory Stalls are counteracted by larger CPU caches, and hyper-threading.

HYPER-THREADING

Hyper-Threading, also known as Hardware multithreading, Simultaneous Multithreading, or Chip multithreading involves a CPU core presenting itself as two cores. In reality there is only one core, capable of switching between two threads of execution in case of a memory stall, i.e. if thread 1 stalls the CPU starts executing thread 2. In some cases even the OS may not be aware of hyper-threading, which may cause problems on a multi core system, where a scheduler may in theory schedule two processes to run on logical cores belonging to the same physical core, leaving one core running two processes and the other core idle. This problem is solved by making the OS aware of hyper-threading.

Vika 8

ILLUSTRATE A RACE CONDITION AND DESCRIBE THE CRITICAL-SECTION PROBLEM.

A race condition occurs when two or more threads access shared data concurrently, and at least one of them modifies the data, causing unexpected behavior due to the unpredictable order of execution.

The critical-section problem refers to the challenge of ensuring that when a thread is executing in its critical section (where shared data is accessed), no other thread can enter its critical section. This helps prevent race conditions and maintain data consistency.

PRESENT SOFTWARE, HARDWARE AND HIGHER-LEVEL SOLUTIONS TO THE CRITICAL-SECTION PROBLEM:

- Software: Peterson's algorithm,
- Hardware: Atomic instructions, including spinlocks,
- Higher-level: Semaphores, monitors and condition variables, message passing.
- Java API for semaphores and monitors with conditions.

PETERSON'S ALGORITHM,

The Peterson algorithm is a concurrent programming algorithm for mutual exclusion that allows two processes to share a single-use resource without conflict, using only shared memory for communication. It was formulated by Gary L. Peterson in 1981. The algorithm uses two variables, flag and turn, to indicate whether a process is ready to enter the critical section or not.

ATOMIC INSTRUCTIONS, INCLUDING SPINLOCKS,

An operation or instruction (or a region of several instructions) is atomic, if it appears to the rest of the system to occur instantaneously. Atomicity is a guarantee of isolation from concurrent processes. Additionally, atomic operations commonly have a succeed-or-fail definition — they either successfully change the state of the system, or have no apparent effect.

SEMAPHORES,

A semaphore can be considered as a special type of variable (or as a special class in case of object-orientation). It consists of:

- A value,
- A queue of threads waiting for the value to be greater than zero.
- Operations that can be applied to it:
 - Wait: decrement the value, and if it is negative, block the thread until it becomes positive again,
- Signal: increment the value, and if it was negative, unblock one of the threads waiting for the semaphore.
- Initialize: set the value to a given number.

MUTUAL EXCLUSION

```
sem_condition.init(1);

// Process A:
sem_condition.wait();
//<critical section>
sem_condition.signal();

// Process B:
sem_condition.wait();
//<critical section>
sem_condition.signal();
```

MONITORS AND CONDITION VARIABLES,

Monitors: A monitor is a programming construct that enforces mutual exclusion by allowing only one thread to execute within its critical section at a time. It is an object that encapsulates both data (shared resources) and the methods (functions) that operate on that data. Monitors ensure that the methods are executed atomically, preventing race conditions.

Condition variables: Condition variables are used in conjunction with monitors to manage threads that must wait for specific conditions to be met before they can proceed. They allow threads to wait for a certain condition within the monitor and be notified when the condition is met. Condition variables provide an efficient way to manage threads that need to wait and resume execution based on specific conditions.

MESSAGE PASSING.

Message passing is an alternative approach to shared memory for exchanging data between threads or processes. Instead of using shared memory regions, which require synchronization mechanisms like semaphores, monitors, or condition variables, message passing relies on explicit communication via messages sent between the cooperating entities.

Message passing systems can be implemented using various communication channels like pipes, sockets, message queues, or higher-level APIs provided by programming languages or libraries. These systems can be synchronous (blocking) or asynchronous (non-blocking) based on how communication is managed.

Encapsulation: Message passing promotes the encapsulation of data, as threads or processes do not directly access shared memory. Instead, they communicate by exchanging messages containing the required data. **Synchronization:** With message passing, synchronization is often implicit, as sending and receiving messages may involve blocking or non-blocking behavior, depending on the

implementation. When a sender is blocked until the receiver accepts the message, it enforces a natural synchronization point.

Scalability: Message passing can be more scalable than shared memory, especially in distributed systems, as it does not rely on a single shared memory region. This makes it more suitable for communication across different systems or networks.

Ease of reasoning: Since message passing avoids direct manipulation of shared memory, it can be easier to reason about the correctness and safety of concurrent programs. However, it may require more explicit communication and handling of message passing events.

JAVA API FOR SEMAPHORES AND MONITORS WITH CONDITIONS.

```
import java.util.concurrent.Semaphore;

try {
    semaphore.acquire();
    // start of critical section
    // ...
    // end of critical section
    sem.release();
} catch (InterruptedException e) {
    // i.e. someone called interrupt() while we were waiting in
    // sem.acquire() call
}
```

Monitor construct is based on the encapsulation of data within an object, and the use of the synchronized keyword to ensure that only one thread can access the object at a time. A method can be declared as synchronized, or a block of code can be synchronized. And if a thread is executing a synchronized method or block, no other thread can execute any other synchronized method or block on the same object.

PRESENT CLASSICAL SYNCHRONISATION PROBLEMS.

BOUNDED-BUFFER PROBLEM,

Same as producer-consumer problem,

READERS-WRITERS PROBLEM,

There are at least three variations of the problems, which deal with situations in which many concurrent threads of execution try to access the same shared resource at one time. Some threads may read and some may write, with the constraint that no thread may access the shared resource for either reading or writing while another thread is in the act of writing to it. (In particular, we want to prevent more than one thread modifying the shared resource simultaneously and allow for two or more readers to access the shared resource at the same time).

DINING PHILOSOPHERS PROBLEM.

Five philosophers dine together at the same table. Each philosopher has their own place at the table. There is a fork between each plate. The dish served is a kind of spaghetti which has to be eaten with two forks. Each philosopher can only alternately think and eat. Moreover, a philosopher can only eat their spaghetti when they have both a left and right fork. Thus two forks will only be available when their two nearest neighbors are thinking, not eating. After an individual philosopher finishes eating, they will put down both forks.

Deadlocks: Imagine each philosopher picks up their left fork. Then each philosopher will wait forever for their right fork. This is a deadlock.

Starvation: Imagine that each philosopher always picks up their left fork first. Then each philosopher will wait forever for their right fork. This is starvation.

Solution (not 100% correct but enough to get the feeling for how it is done):

```
class Philosopher extends Thread {
    private final Semaphore leftFork;
    private final Semaphore rightFork;

    Philosopher(Semaphore leftFork, Semaphore rightFork) {
        this.leftFork = leftFork;
        this.rightFork = rightFork;
    }

    public void run() {
        while (true) {
            think();
            pickUpForks();
            eat();
            putDownForks();
        }
    }

    private void think() {
        // Philosopher is thinking
    }

    private void pickUpForks() {
        try {
            leftFork.acquire();
            rightFork.acquire();
        } catch (InterruptedException e) {
            e.printStackTrace();
        }
    }

    private void eat() {
        // Philosopher is eating
    }

    private void putDownForks() {
        leftFork.release();
        rightFork.release();
    }
}

public class DiningPhilosophers {
    public static void main(String[] args) {
        int numberOfPhilosophers = 5;
        Semaphore[] forks = new Semaphore[numberOfPhilosophers];
        Philosopher[] philosophers = new Philosopher[numberOfPhilosophers];

        for (int i = 0; i < numberOfPhilosophers; i++) {
            forks[i] = new Semaphore(1); // Each fork is a semaphore with 1 permit
        }

        for (int i = 0; i < numberOfPhilosophers; i++) {
            Semaphore leftFork = forks[i];
            Semaphore rightFork = forks[(i + 1) % numberOfPhilosophers];
            philosophers[i] = new Philosopher(leftFork, rightFork);
            philosophers[i].start();
        }
    }
}
```

SLEEPING BARBER PROBLEM.

Imagine a hypothetical barbershop with one barber, one barber chair, and a waiting room with n chairs (n may be 0) for waiting customers. The following rules apply:

- If there are no customers, the barber falls asleep in the chair
- If a customer arrives and the barber is asleep, the customer wakes up the barber.
- When a customer arrives while the barber is cutting someone else's hair, he sits down in one of the chairs in the waiting room.
- If there are no empty chairs, the customer leaves.
- When the barber finishes cutting a customer's hair, they dismiss the customer and return to the barber chair to sleep if there are no other customers waiting.

```
Semaphore ME = new Semaphore(1); // Mutex for the waiting room
Semaphore barberSleep = new Semaphore(0); // Initially asleep
Semaphore barberChair = new Semaphore(0); // Mutex for the barber chair
int numberOfFreeWaitRoomSeats = N; // Number of free seats in the waiting room

// Barber:
void Barber () {
    while (true) {
        barberSleep.wait(); // Try to sleep
        ME.wait(); // Enter the waiting room
        numberOfFreeWaitRoomSeats++; // One chair becomes free
        barberChair.signal(); // Invite customer into the chair
        cutHairOfCustomerOnChair(); // Cut hair
        ME.signal(); // Release the waiting room
    }
}

// Customer:
void Customer() {
    ME.wait(); // Enter the waiting room
    if (numberOfFreeWaitRoomSeats > 0) {
        numberOfFreeWaitRoomSeats--; // Occupy a chair
        barberSleep.signal(); // Wake up the barber if needed
        ME.signal(); // Release the waiting room
        barberChair.wait(); // Wait until invited
        goToBarberChairGetHairCutLeave(); // Get haircut
    } else {
        ME.signal(); // Release the waiting room
        leaveWithoutHaircut(); // No free chairs
    }
}
```